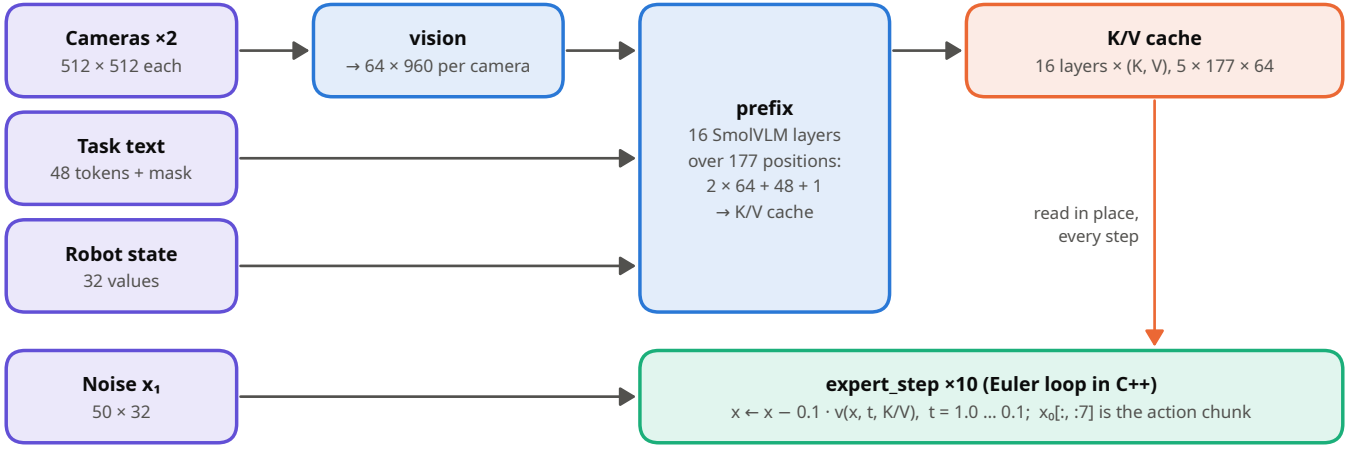
Cameras ×2
512 × 512 each
vision
→ 64 × 960 per camera
Task text
48 tokens + mask
Robot state
32 values
prefix
16 SmolVLM layers
over 177 positions:
2 × 64 + 48 + 1
→ K/V cache
K/V cache
16 layers × (K, V), 5 × 177 × 64
read in place,
every step
Noise x₁
50 × 32
expert_step ×10 (Euler loop in C++)
x ← x − 0.1 · v(x, t, K/V), t = 1.0 ... 0.1; x₀[:, :7] is the action chunk